

Problem Set 6

Prasanna Parasurama
(Dated: May 10, 2019)

I. PROBLEM 1

This exercise asks you to implement and assess the performance of the bootstrap for the linear regression model. Suppose you have the linear regression model:

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i$$

where,

- $x_i \sim U[0, 2]$
- $\epsilon_i | x_i \sim U[-1, 1]$
- $\beta_0 = \beta_1 = 1$

We ask you to answer the following questions:

- a. Write a code that generates i.i.d. samples of sizes $n = 10, 50, 200$ from that distribution, computes (1) the least squares estimator for β , (2) the t-ratio for the least squares coefficient β_1 , $t_n = \frac{\hat{\beta}_{1,LS} - 1}{s.e.(\hat{\beta}_{1,LS})}$, and (3) the least square residuals $\hat{\epsilon}_i = y_i - \hat{\beta}_{0,LS} - \hat{\beta}_{1,LS}x_i$

```
In [1]: import numpy as np
import pandas as pd
import statsmodels.api as sm
import matplotlib.pyplot as plt
from datetime import datetime
%matplotlib inline

In [2]: class BootstrapSimulator:
    def __init__(self, n):
        self.n = n # sample size

        self.df = pd.DataFrame()

        self.ols = None
        self.b0_ols = None
        self.b1_ols = None
        self.t_ols = None

    def draw_samples(self):
        n = self.n

        self.df['x'] = np.random.uniform(0, 2, n)
        self.df['e'] = np.random.uniform(-1, 1, n)
        self.df['y'] = 1 + self.df['x'] + self.df['e']
        self.df['intercept'] = np.ones(n)
        return self.df

    @staticmethod
    def get_ols_params(fitted_model):
```

```

        b_hat = fitted_model.params['x']
        b_se = fitted_model.bse['x']
        t = (b_hat - 1) / b_se
        return b_hat, t

def fit_ols(self, y="y"):
    # fit model
    self.ols = sm.OLS(self.df[y], self.df[['x', 'intercept']]).fit()

    self.b1_ols, self.t_ols = self.get_ols_params(self.ols)
    self.b0_ols = self.ols.params['intercept']

    # residuals
    self.df['e_hat'] = self.df[y] - self.b0_ols - self.b1_ols * self.df['x']

    return self.b1_ols, self.t_ols

def sample_residuals(self):
    n = len(self.df)
    return np.random.choice(self.df['e_hat'], n, replace=True)

@staticmethod
def get_bootstrap_stats(params):
    return {'ci': (np.percentile(params, 2.5), np.percentile(params, 97.5)),
            'mean': np.mean(params)}

def residual_bootstrap(self, n_reps=200):

    b_bootstrap = list()
    t_bootstrap = list()

    stats = dict()

    for _ in range(n_reps):
        e_bootstrap = self.sample_residuals()
        self.df['y_bootstrap'] = self.b0_ols + self.b1_ols * self.df['x'] +
e_bootstrap
        ols = sm.OLS(self.df["y_bootstrap"], self.df[['x', 'intercept']]).fit()
        b, t = self.get_ols_params(ols)
        b_bootstrap.append(b)
        t_bootstrap.append(t)

    stats['b'] = self.get_bootstrap_stats(b_bootstrap)
    stats['t'] = self.get_bootstrap_stats(t_bootstrap)

    return stats

def parametric_bootstrap(self, n_reps=200):

    b_bootstrap = list()
    t_bootstrap = list()

    stats = dict()

    for _ in range(n_reps):
        self.df['e'] = np.random.uniform(-1, 1, self.n)
        self.df['y'] = self.b0_ols + self.b1_ols * self.df['x'] + self.df['e']

```

```

ols = sm.OLS(self.df['y'], self.df[['x', 'intercept']]).fit()

b, t = self.get_ols_params(ols)
b_bootstrap.append(b)
t_bootstrap.append(t)
stats['b'] = self.get_bootstrap_stats(b_bootstrap)
stats['t'] = self.get_bootstrap_stats(t_bootstrap)

return stats

```

See `draw_samples` and `fit_ols` methods.

- b. Write a code for drawing n times at random from the discrete uniform distribution over the estimated residuals $\hat{\epsilon}_1, \dots, \hat{\epsilon}_n$ (i.e. with replacement).

See `sample_residuals` method in `BootstrapSimulator` class.

- c. Use your code from parts (a) and (b) to implement the residual bootstrap - assuming that ϵ_i and x_i are independent - to estimate the 95th percentiles of the respective distributions of $\hat{\beta}_{1,LS}$ and t_n

```

In [3]: bs_10 = BootstrapSimulator(n=10)
        bs_50 = BootstrapSimulator(n=50)
        bs_200 = BootstrapSimulator(n=200)

        bs_10.draw_samples()
        bs_10.fit_ols()

        bs_50.draw_samples()
        bs_50.fit_ols()

        bs_200.draw_samples()
        bs_200.fit_ols()

Out[3]: (0.9850883905439032, -0.19459470136382742)

In [4]: # Residual bootstrap
        res_10 = bs_10.residual_bootstrap()
        res_50 = bs_50.residual_bootstrap()
        res_200 = bs_200.residual_bootstrap()

In [5]: df_res = pd.DataFrame([res_10, res_50, res_200])

        print('b estimates')
        print(df_res['b'].apply(pd.Series))

        print('t estimates')
        print(df_res['t'].apply(pd.Series))

```

b estimates

	ci	mean
0	(0.1252409779586813, 0.8744245281189442)	0.497854
1	(0.7564703092768035, 1.3095184865086547)	1.024620
2	(0.8415242880163574, 1.1121380817466304)	0.975943

t estimates

	ci	mean
0	(-4.9693357957442394, -0.5040697930336879)	-2.357089
1	(-1.8305105346400985, 2.240661045221299)	0.177186
2	(-2.016877673129661, 1.467179807545017)	-0.316971

- d. Repeat part (a) for sample size $n = 10, 50, 200$ with 200 replications, where you keep the initial draws of x_1, \dots, x_n from part (a) and only generate new residuals from their conditional distribution. Compute $\hat{\beta}_{1,LS}$ and the statistic t_n using 200 independent samples of size n . Use your results to compute a simulated estimate for the 95th percentiles of the respective sampling distributions for $\hat{\beta}_{1,LS}$ and t_n .

```
In [6]: # parametric bootstrap
par_10 = bs_10.parametric_bootstrap()
par_50 = bs_50.parametric_bootstrap()
par_200 = bs_200.parametric_bootstrap()
```

```
In [7]: df_par = pd.DataFrame([par_10, par_50, par_200])
```

```
print('b estimates')
print(df_par['b'].apply(pd.Series))

print('t estimates')
print(df_par['t'].apply(pd.Series))
```

```
b estimates
                                     ci      mean
0  (-0.07770851234812026, 1.0908798535407562)  0.481773
1    (0.7336699250353068, 1.2440893219167681)  1.014423
2    (0.8299080953533149, 1.1223787253665154)  0.985627
t estimates
                                     ci      mean
0  (-4.249443952035807, 0.3108292477213074) -1.776718
1  (-2.0759239890186234, 1.7110379456127023)  0.091293
2  (-2.3983362460727142, 1.6721542243509722) -0.196839
```

- e. Compare your results from (c) and (d). What do you conclude about the performance of the bootstrap? How does it compare to the 95th percentile of the asymptotic distribution of t_n ?

Both residual (c) and parametric (d) bootstrap give similar results for *beta* estimates. For *t* estimates, they are different for $n=2$, but tend to converge for higher n .

```
In [8]: def get_asymptotic_ci(bs_sim, stat):
        if stat == "b":
            return (bs_sim.ols.params['x'] - 2*bs_sim.ols.bse['x'], bs_sim.ols.params['x'] +
                    2*bs_sim.ols.bse['x'])
        elif stat == "t":
            return (((bs_sim.ols.params['x'] - 2*bs_sim.ols.bse['x']) -
                    1)/bs_sim.ols.bse['x'], ((bs_sim.ols.params['x'] + 2*bs_sim.ols.bse['x']) -
                    1)/bs_sim.ols.bse['x'])
```

For, $\hat{\beta}_1$, note that the bootstrap CI are fairly close to asymptotic distribution CI:

```
In [9]: print(get_asymptotic_ci(bs_10, stat="b"))
        print(get_asymptotic_ci(bs_50, stat="b"))
        print(get_asymptotic_ci(bs_200, stat="b"))
```

```
(0.0012296909768160225, 1.0016978646647163)
(0.7340372030038689, 1.3007819814839865)
(0.831830266334066, 1.1383465147537404)
```

To get asymptotic distribution of t_n , we would have to rely on the delta method. But since β_1 and $se(\beta_1)$ are independent, we can do a simple substitution. The bootstrap estimates converge for larger n .

```
In [10]: print(get_asymptotic_ci(bs_10, stat="t"))
         print(get_asymptotic_ci(bs_50, stat="t"))
         print(get_asymptotic_ci(bs_200, stat="t"))
```

```
(-3.9932117194354806, 0.00678828056451893)
(-1.877125698162645, 2.122874301837354)
(-2.1945947013638274, 1.8054052986361726)
```

II. PROBLEM 2

This exercise will walk you through a prediction task. I have downloaded data from a peer-to-peer lending platform, Lending Club. The dataset you will work with is: **`lending_club_07_to_11_cleaned.csv`**. Lending Club provides detailed characteristic information regarding loans, both information on the borrower, as well as, the loan itself. Your goal will be to build a model to predict the outcome of a loan, i.e. whether an individual paid off a loan or did not pay off a loan. In our case, a good outcome is if the loan is fully paid off, a bad outcome is if the loan is charged off.

The target variable for the analysis is **`loan_status`**, where:

$$loan_status = \begin{cases} 1 & \text{if loan is paid off} \\ 0 & \text{if loan is not paid off} \end{cases}$$

```
In [11]: import pandas as pd
```

```
df = pd.read_csv("data/lending_club_07_to_11_cleaned.csv")
```

- a. This is going to be a more DIY style exercise, provide a list of the variables you plan to use for the analysis. Give a short discussion for why you excluded other variables.

```
In [12]: vars_to_include = ["addr_state",
                           "annual_inc",
                           "delinq_2yrs",
                           "dti",
                           "emp_length",
                           "funded_amnt",
                           "home_ownership",
                           "inq_last_6mths",
                           "installment",
                           "int_rate",
                           'loan_amnt',
                           'open_acc',
                           'pub_rec',
                           'pub_rec_bankruptcies',
                           'purpose',
                           'revol_bal',
                           'revol_util',
                           'sub_grade',
                           'term',
                           'total_acc',
                           'total_pymnt',
                           'total_rec_int',
                           'total_rec_late_fee',
                           'total_rec_prncp',
                           'verification_status',
                           'zip_code']
```

```

In [13]: # convert to datetime
df['issue_d'] = pd.to_datetime(df['issue_d'], format='%b-%y')
df['earliest_cr_line'] = pd.to_datetime(df['earliest_cr_line'], format='%b-%y')
df['last_pymnt_d'] = pd.to_datetime(df['last_pymnt_d'], format='%b-%y')
df['last_credit_pull_d'] = pd.to_datetime(df['last_credit_pull_d'], format='%b-%y')

In [14]: # create new features
df['days_since_first_cr_line'] = (df['issue_d'] - df['earliest_cr_line']).apply(lambda x: x.days)
df['days_since_last_pymnt'] = (datetime.today() - df['last_pymnt_d']).apply(lambda x: x.days)
df['days_since_last_cr_pull'] = (datetime.today() - df['last_credit_pull_d']).apply(lambda x: x.days)

In [15]: vars_to_include.extend(['days_since_first_cr_line', 'days_since_last_pymnt',
                                'days_since_last_cr_pull'])

In [16]: df[vars_to_include].dtypes

Out[16]: addr_state          object
annual_inc          float64
delinq_2yrs          int64
dti                 float64
emp_length          object
funded_amnt          int64
home_ownership       object
inq_last_6mths       int64
installment          float64
int_rate             object
loan_amnt            int64
open_acc             int64
pub_rec              int64
pub_rec_bankruptcies float64
purpose              object
revol_bal            int64
revol_util           object
sub_grade            object
term                 object
total_acc            int64
total_pymnt          float64
total_rec_int         float64
total_rec_late_fee    float64
total_rec_prncp       float64
verification_status  object
zip_code             object
days_since_first_cr_line int64
days_since_last_pymnt float64
days_since_last_cr_pull float64
dtype: object

In [17]: # convert to proper data types
df['int_rate'] = df['int_rate'].apply(lambda x: float(x[:-1]) if x is not np.nan else x)
df['revol_util'] = df['revol_util'].apply(lambda x: float(x[:-1]) if x is not np.nan else x)

```

The following variables are omitted, because there is no variance in the data: acc_now_delinq, chargeoff_within_12_mths, collections_12_mths_ex_med, delinq_amnt, disbursement method,

hardship_flag, initial_list_status, out_prncp, out_prncp_inv, policy_code, pymnt_plan, tax_liens.

The following variables are omitted because it is not in the data dictionary, and I can't infer what they mean: collection_recovery_fee, debt_settlement_flag, funded_amnt_inv, total_pymnt_inv, recoveries.

grade is fully capture in sub_grade.

Instead of including, issue_d, earliest_cr_line, last_credit_pull_d, last_pymnt_d, I included 3 transformed features: days_since_first_cr_line, days_since_last_pymnt, days_since_last_pymnt

- b. Regularization is an important step when using an machine learning algorithm, regularize the variables that you have included. Briefly, why is regularization important?

Regularization centers and scales the data. It helps with gradient descent convergence as the gradient surface won't be skewed in any particular dimension

```
In [18]: data = df[vars_to_include+['loan_status']]
        data = data.dropna().reset_index(drop=True)
```

```
X = data[vars_to_include]
```

```
In [19]: from sklearn.preprocessing import StandardScaler
```

```
data = df[vars_to_include+['loan_status']]
data = data.dropna().reset_index(drop=True)
```

```
X = data[vars_to_include]
y = data['loan_status']
# get numeric column types
```

```
numeric_cols = [k for k,v in X.dtypes.to_dict().items() if v in [int, float]]
print(numeric_cols)
scaler = StandardScaler()
```

```
X[numeric_cols] = scaler.fit_transform(X[numeric_cols])
```

```
['annual_inc', 'delinq_2yrs', 'dti', 'funded_amnt', 'inq_last_6mths', 'installment',
'int_rate', 'loan_amnt', 'open_acc', 'pub_rec', 'pub_rec_bankruptcies', 'revol_bal',
'revol_util', 'total_acc', 'total_pymnt', 'total_rec_int', 'total_rec_late_fee',
'total_rec_prncp', 'days_since_first_cr_line', 'days_since_last_pymnt',
'days_since_last_cr_pull']
```

/Users/parasu/anaconda3/lib/python3.6/site-packages/ipykernel_launcher.py:14:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

/Users/parasu/anaconda3/lib/python3.6/site-packages/pandas/core/indexing.py:543:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

```
self.obj[item] = s
```

- c. Provide a simple correlational table to give you a sense of the relationship between your covariates. Do you notice any interesting patterns?

In [20]: `import seaborn as sns`

```
cor_table = X[numeric_cols].corr()
```

```
cor_table
```

```
Out[20]:
```

	annual_inc	delinq_2yrs	dti	funded_amnt	\
annual_inc	1.000000	0.022445	-0.124681	0.268287	
delinq_2yrs	0.022445	1.000000	-0.036118	-0.034389	
dti	-0.124681	-0.036118	1.000000	0.065301	
funded_amnt	0.268287	-0.034389	0.065301	1.000000	
inq_last_6mths	0.033472	0.007069	-0.000236	0.007624	
installment	0.272153	-0.022195	0.053239	0.955391	
int_rate	0.052448	0.158380	0.108273	0.313052	
loan_amnt	0.272296	-0.033730	0.065034	0.981301	
open_acc	0.157171	0.011986	0.290514	0.172354	
pub_rec	-0.015147	0.009970	-0.005554	-0.048898	
pub_rec_bankruptcies	-0.012931	0.004945	0.005774	-0.034133	
revol_bal	0.279054	-0.056674	0.228631	0.313264	
revol_util	0.017173	-0.044628	0.276073	0.069696	
total_acc	0.237203	0.068196	0.229083	0.249831	
total_pymnt	0.257571	-0.023685	0.064530	0.903647	
total_rec_int	0.185332	0.022342	0.104411	0.736995	
total_rec_late_fee	0.008124	0.033372	-0.010295	0.052480	
total_rec_prncp	0.259872	-0.039953	0.041901	0.874004	
days_since_first_cr_line	0.180313	0.064830	0.050192	0.193477	
days_since_last_pymnt	-0.015844	0.005000	-0.097698	-0.146958	
days_since_last_cr_pull	0.010625	-0.007269	-0.119076	-0.042719	

	inq_last_6mths	installment	int_rate	loan_amnt	\
annual_inc	0.033472	0.272153	0.052448	0.272296	
delinq_2yrs	0.007069	-0.022195	0.158380	-0.033730	
dti	-0.000236	0.053239	0.108273	0.065034	
funded_amnt	0.007624	0.955391	0.313052	0.981301	
inq_last_6mths	1.000000	0.007526	0.134738	0.007576	
installment	0.007526	1.000000	0.282471	0.929132	
int_rate	0.134738	0.282471	1.000000	0.309604	
loan_amnt	0.007576	0.929132	0.309604	1.000000	
open_acc	0.092757	0.169679	0.013577	0.173914	
pub_rec	0.024603	-0.043192	0.097092	-0.048105	
pub_rec_bankruptcies	0.015048	-0.029963	0.082078	-0.032866	
revol_bal	-0.025706	0.315640	0.100515	0.320136	
revol_util	-0.068598	0.095136	0.465695	0.065284	
total_acc	0.113206	0.230644	-0.042612	0.255719	
total_pymnt	-0.010114	0.854238	0.313651	0.886746	
total_rec_int	0.023080	0.632747	0.531746	0.729106	
total_rec_late_fee	0.029547	0.059798	0.100839	0.050447	
total_rec_prncp	-0.023412	0.851132	0.193289	0.855311	
days_since_first_cr_line	0.006994	0.169246	-0.093930	0.199594	
days_since_last_pymnt	0.073552	-0.043235	-0.102115	-0.146459	
days_since_last_cr_pull	-0.025196	-0.019885	-0.119939	-0.039591	

	open_acc	pub_rec	...	revol_bal	revol_util	\
annual_inc	0.157171	-0.015147	...	0.279054	0.017173	

delinq_2yrs	0.011986	0.009970	...	-0.056674	-0.044628
dti	0.290514	-0.005554	...	0.228631	0.276073
funded_amnt	0.172354	-0.048898	...	0.313264	0.069696
inq_last_6mths	0.092757	0.024603	...	-0.025706	-0.068598
installment	0.169679	-0.043192	...	0.315640	0.095136
int_rate	0.013577	0.097092	...	0.100515	0.465695
loan_amnt	0.173914	-0.048105	...	0.320136	0.065284
open_acc	1.000000	0.003568	...	0.286259	-0.087536
pub_rec	0.003568	1.000000	...	-0.059926	0.057715
pub_rec_bankruptcies	0.009457	0.841247	...	-0.047077	0.060246
revol_bal	0.286259	-0.059926	...	1.000000	0.305062
revol_util	-0.087536	0.057715	...	0.305062	1.000000
total_acc	0.684804	-0.020384	...	0.312378	-0.068831
total_pymnt	0.158537	-0.050642	...	0.295432	0.080870
total_rec_int	0.122305	-0.005121	...	0.244697	0.195013
total_rec_late_fee	-0.016989	-0.006129	...	0.006222	0.040979
total_rec_prncp	0.156791	-0.062609	...	0.284513	0.026387
days_since_first_cr_line	0.208382	0.035276	...	0.245795	-0.025639
days_since_last_pymnt	-0.027757	0.013952	...	-0.034163	-0.090850
days_since_last_cr_pull	-0.074567	-0.052309	...	-0.045488	-0.145986

	total_acc	total_pymnt	total_rec_int	\
annual_inc	0.237203	0.257571	0.185332	
delinq_2yrs	0.068196	-0.023685	0.022342	
dti	0.229083	0.064530	0.104411	
funded_amnt	0.249831	0.903647	0.736995	
inq_last_6mths	0.113206	-0.010114	0.023080	
installment	0.230644	0.854238	0.632747	
int_rate	-0.042612	0.313651	0.531746	
loan_amnt	0.255719	0.886746	0.729106	
open_acc	0.684804	0.158537	0.122305	
pub_rec	-0.020384	-0.050642	-0.005121	
pub_rec_bankruptcies	-0.007656	-0.039649	-0.000577	
revol_bal	0.312378	0.295432	0.244697	
revol_util	-0.068831	0.080870	0.195013	
total_acc	1.000000	0.222062	0.145692	
total_pymnt	0.222062	1.000000	0.836404	
total_rec_int	0.145692	0.836404	1.000000	
total_rec_late_fee	-0.020286	0.021781	0.082799	
total_rec_prncp	0.228903	0.971794	0.695599	
days_since_first_cr_line	0.358805	0.179525	0.123714	
days_since_last_pymnt	-0.003767	-0.317523	-0.474569	
days_since_last_cr_pull	-0.066941	-0.023946	-0.131297	

	total_rec_late_fee	total_rec_prncp	\
annual_inc	0.008124	0.259872	
delinq_2yrs	0.033372	-0.039953	
dti	-0.010295	0.041901	
funded_amnt	0.052480	0.874004	
inq_last_6mths	0.029547	-0.023412	
installment	0.059798	0.851132	
int_rate	0.100839	0.193289	
loan_amnt	0.050447	0.855311	
open_acc	-0.016989	0.156791	
pub_rec	-0.006129	-0.062609	
pub_rec_bankruptcies	-0.005562	-0.050157	

revol_bal	0.006222	0.284513
revol_util	0.040979	0.026387
total_acc	-0.020286	0.228903
total_pymnt	0.021781	0.971794
total_rec_int	0.082799	0.695599
total_rec_late_fee	1.000000	-0.013677
total_rec_prncp	-0.013677	1.000000
days_since_first_cr_line	-0.010749	0.183498
days_since_last_pymnt	0.008729	-0.237988
days_since_last_cr_pull	-0.026863	0.023259

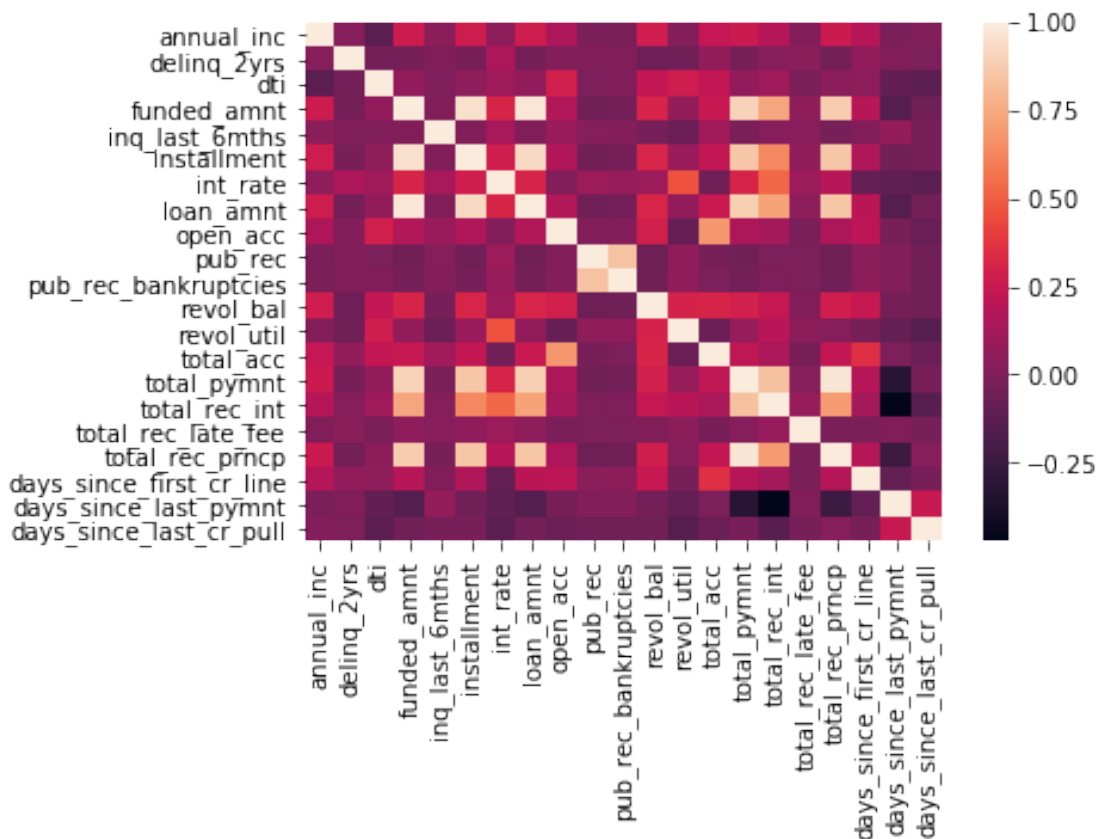
	days_since_first_cr_line	days_since_last_pymnt \
annual_inc	0.180313	-0.015844
delinq_2yrs	0.064830	0.005000
dti	0.050192	-0.097698
funded_amnt	0.193477	-0.146958
inq_last_6mths	0.006994	0.073552
installment	0.169246	-0.043235
int_rate	-0.093930	-0.102115
loan_amnt	0.199594	-0.146459
open_acc	0.208382	-0.027757
pub_rec	0.035276	0.013952
pub_rec_bankruptcies	0.044421	0.012104
revol_bal	0.245795	-0.034163
revol_util	-0.025639	-0.090850
total_acc	0.358805	-0.003767
total_pymnt	0.179525	-0.317523
total_rec_int	0.123714	-0.474569
total_rec_late_fee	-0.010749	0.008729
total_rec_prncp	0.183498	-0.237988
days_since_first_cr_line	1.000000	-0.088809
days_since_last_pymnt	-0.088809	1.000000
days_since_last_cr_pull	-0.024030	0.246529

	days_since_last_cr_pull
annual_inc	0.010625
delinq_2yrs	-0.007269
dti	-0.119076
funded_amnt	-0.042719
inq_last_6mths	-0.025196
installment	-0.019885
int_rate	-0.119939
loan_amnt	-0.039591
open_acc	-0.074567
pub_rec	-0.052309
pub_rec_bankruptcies	-0.046864
revol_bal	-0.045488
revol_util	-0.145986
total_acc	-0.066941
total_pymnt	-0.023946
total_rec_int	-0.131297
total_rec_late_fee	-0.026863
total_rec_prncp	0.023259
days_since_first_cr_line	-0.024030
days_since_last_pymnt	0.246529
days_since_last_cr_pull	1.000000

[21 rows x 21 columns]

```
In [21]: sns.heatmap(X[numeric_cols].corr())
```

```
Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x12e128898>
```



```
In [22]: X[['total_rec_prncp', 'total_pymnt']].corr()
```

```
Out[22]:
```

	total_rec_prncp	total_pymnt
total_rec_prncp	1.000000	0.971794
total_pymnt	0.971794	1.000000

There are some almost perfect correlation, for example total_rec_prncp, which we can take out.

```
In [23]: X = X.drop('total_rec_prncp', axis=1)
```

d. Split the dataset into a single test and training set, a simple rule of thumb is an 40/60 split. How did you build these two sets?

```
In [24]: from sklearn.model_selection import train_test_split
```

```
# first convert categorical cols to dummies
cat_cols = [k for k,v in X.dtypes.to_dict().items() if v in [object]]
dummy = pd.get_dummies(X[cat_cols])
X = pd.concat([X, dummy], axis=1)
X = X.drop(cat_cols, axis=1)
```

```
In [25]: X.shape
```

```
Out[25]: (37901, 947)
```

```
In [26]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4)
```

- e. Using your training set, run a logistic regression, a random forest, and a gradient boosted random forest. To show your results, present both a measure of misclassification error, accuracy and a confusion matrix.

```
In [27]: from sklearn.linear_model import LogisticRegression
         from sklearn.ensemble import RandomForestClassifier
         from sklearn.ensemble import GradientBoostingClassifier
         from sklearn.metrics import accuracy_score, confusion_matrix
```

```
         rf = RandomForestClassifier()
         logit = LogisticRegression()
         rfgb = GradientBoostingClassifier()
```

```
/Users/parasu/anaconda3/lib/python3.6/site-
packages/sklearn/ensemble/weight_boosting.py:29: DeprecationWarning:
numpy.core.umath_tests is an internal NumPy module and should not be imported. It will
be removed in a future NumPy release.
    from numpy.core.umath_tests import inner1d
```

```
In [28]: logit.fit(X_train, y_train)
         rf.fit(X_train, y_train)
         rfgb.fit(X_train, y_train)
```

```
Out[28]: GradientBoostingClassifier(criterion='friedman_mse', init=None,
                                     learning_rate=0.1, loss='deviance', max_depth=3,
                                     max_features=None, max_leaf_nodes=None,
                                     min_impurity_decrease=0.0, min_impurity_split=None,
                                     min_samples_leaf=1, min_samples_split=2,
                                     min_weight_fraction_leaf=0.0, n_estimators=100,
                                     presort='auto', random_state=None, subsample=1.0, verbose=0,
                                     warm_start=False)
```

```
In [29]: def get_metrics(model):
         print(model)
         print("accuracy")
         print(accuracy_score(y_pred=model.predict(X_test), y_true=y_test))
         print("Confusion matrix")
         print(confusion_matrix(y_pred=model.predict(X_test), y_true=y_test))
         print('\n\n')
```

```
In [30]: for m in (rf, logit, rfgb):
         get_metrics(m)
```

```
RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                        max_depth=None, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False)
```

```
accuracy
```

```
0.9389881933909373
```

```
Confusion matrix
```

```
[[ 1251   840]
 [    85 12985]]
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                    verbose=0, warm_start=False)
```

```
accuracy
```

```
0.9827847767297672
```

```
Confusion matrix
```

```
[[ 1832   259]
 [    2 13068]]
```

```
GradientBoostingClassifier(criterion='friedman_mse', init=None,
                           learning_rate=0.1, loss='deviance', max_depth=3,
                           max_features=None, max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=100,
                           presort='auto', random_state=None, subsample=1.0, verbose=0,
                           warm_start=False)
```

```
accuracy
```

```
0.9779038322010422
```

```
Confusion matrix
```

```
[[ 1779   312]
 [   23 13047]]
```

The accuracy seems suspiciously high (the baserate for paid off loan is 86%). Maybe there's leakage from one of the included variables.

- f. One easy way to improve model performance is cross-validation. Do a k-fold cross validation, where k=5, using the best performing model from part (e.). Re-report the misclassification error, accuracy and a confusion matrix.

```
In [31]: from sklearn.model_selection import cross_validate, cross_val_score, cross_val_predict
```

```
In [32]: # accuracies
         cross_val_score(logit, X, y, cv=5)
```

```
Out[32]: array([0.98720485, 0.98680913, 0.98562005, 0.98139842, 0.98113208])
```

```
In [33]: confusion_matrix(y_pred=cross_val_predict(logit, X, y, cv=5), y_true=y)
```

```
Out[33]: array([[ 4663,   581],
                 [    9, 32648]])
```